


# Introduction to Artificial Neural Networks (ANN) – Part 1



**By: Dr Ng Gee Wah**  
**ESP 3201**


# Neural Network In Our Brains



**NUS**  
National University  
of Singapore

## A typical biological neuron comprised of:

- A cell body
- Hair-like dendrites (input channels)
- Axon: output cable



A nerve cell ( neuron)

- Many such nerve cells are arranged together in our brain to form a network of nerves.
- They pass electrical impulses i.e the excitation from one neuron to the other.
- The dendrites receive the impulse from the terminal button or synapse of an adjoining neuron.
- Dendrites carry the impulse to the nucleus of the nerve cell which is also called as soma. The electrical impulse is processed and then passed on to the axon.
- The axon then carries the impulse from the soma to the synapse.
- Next, the synapse passes the impulse to dendrites of the second neuron.
- There are many of this connectivity in the human brain.


# Synaptic Transmission of Neurons

Synaptic transmission is the process whereby one neuron communicates with other neurons, at a synapse.

A typical neuron has a cell body (soma), branching processes specialized to receive incoming signals (dendrites), and a single process (axon) that carries electrical signals away from the neuron toward other neurons. Electrical signals carried by axons are action potentials. Axons often have thousands of terminal branches, each ending as synaptic terminal. At the synaptic terminal, the action potential is converted into a chemical message which, in turn, interacts with the recipient neuron.



**NUS**  
National University  
of Singapore




# Artificial and Biological Neural Networks



**NUS**  
National University  
of Singapore

Biological Neural Network	Artificial Neural Network
Neuron	Processor Element (Unit)
Dendrite	Inputs
Cell Body	Activation Function
Axons	Outputs
Synapses	Weights






# A Single Artificial Neuron

## The Perceptron: Forward Propagation



**NUS**  
National University  
of Singapore



$$\hat{y} = g \left( \omega_0 + \sum_{i=1}^m x_i \omega_i \right)$$

Output

Activation function (Nonlinear in general)

Bias


Linear combination of inputs

# A Single Artificial Neuron

## The Perceptron: Forward Propagation



**NUS**  
National University  
of Singapore




Question: Express the forward propagation equation when  $m=1$  and when  $g$  is

- i) A linear function; say  $g$  is a proportional gain value of 1.
- ii) A non linear function; say  $g$  is a sigmoid function.

# The Simplified Perceptron Output Equation




**NUS**  
National University  
of Singapore



$$\hat{y} = g(w_0 + X^T W)$$

# Exercise case 1




We have:  $w_0 = 1$  and  $\mathbf{W} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$

1. This network only have two inputs. What is the input? (answer:  $x_1$  and  $x_2$ )
2. Given the weights value and  $g$  as a linear function of say gain 1, what is the output of the neuron equation?
3. Find the output of the neuron if the input  $[x_1, x_2] = [-1, 2]$ .



## Exercise case 2




We have:  $w_0 = 1$  and  $\mathbf{W} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$

$$\begin{aligned}\hat{y} &= g(w_0 + \mathbf{X}^T \mathbf{W}) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ \hat{y} &= g(1 + 3x_1 - 2x_2)\end{aligned}$$


This is just a line in 2D!

$g$  is a general nonlinear function in most neuron network design and applications

## Exercise case 2




$$\hat{y} = g(1 + 3x_1 - 2x_2)$$




This is a straight line in the space. This line form the decision boundary.

## Exercise case 2



For  $[x_1, x_2] = [-1, 2]$ ; the output lie somewhere in the space.


## Exercise case 2



Assume we have input:  $\mathbf{X} = \begin{bmatrix} -1 \\ 2 \end{bmatrix}$

$$\begin{aligned}\hat{y} &= g(1 + (3 * -1) - (2 * 2)) \\ &= g(-6) \approx 0.002\end{aligned}$$



$$\hat{y} = g(1 + 3x_1 - 2x_2)$$



Assuming  $g$  is the sigmoid function then you get the output of 0.002.



## Exercise case 2






# Revisit: The Perceptron Equation



**NUS**  
National University  
of Singapore

$$\hat{y} = g(w_0 + \mathbf{X}^T \mathbf{W})$$




1. Dot product of the input and the weight  $\mathbf{X}^T \mathbf{W}$
2. Add the bias term  $w_0$
3. Apply the function. In this case the function is a nonlinear function such as sigmoid function

# The Perceptron: Simplified



**NUS**  
National University  
of Singapore




$$z = w_0 + \sum_{j=1}^m x_j w_j$$

1. We can further simplified the equation and diagram as above. The bias term  $w_0$
2. Apply the function  $g(z)$ . In this case the function,  $g$ , is a nonlinear function such as sigmoid function

# Multiple Output Perceptron




**NUS**  
National University  
of Singapore



$$z_i = w_{0,i} + \sum_{j=1}^m x_j w_{j,i}$$

I will cover this again at Part 2 of this lecture on multiple layers and multiple output.

# A single hidden layer ANN

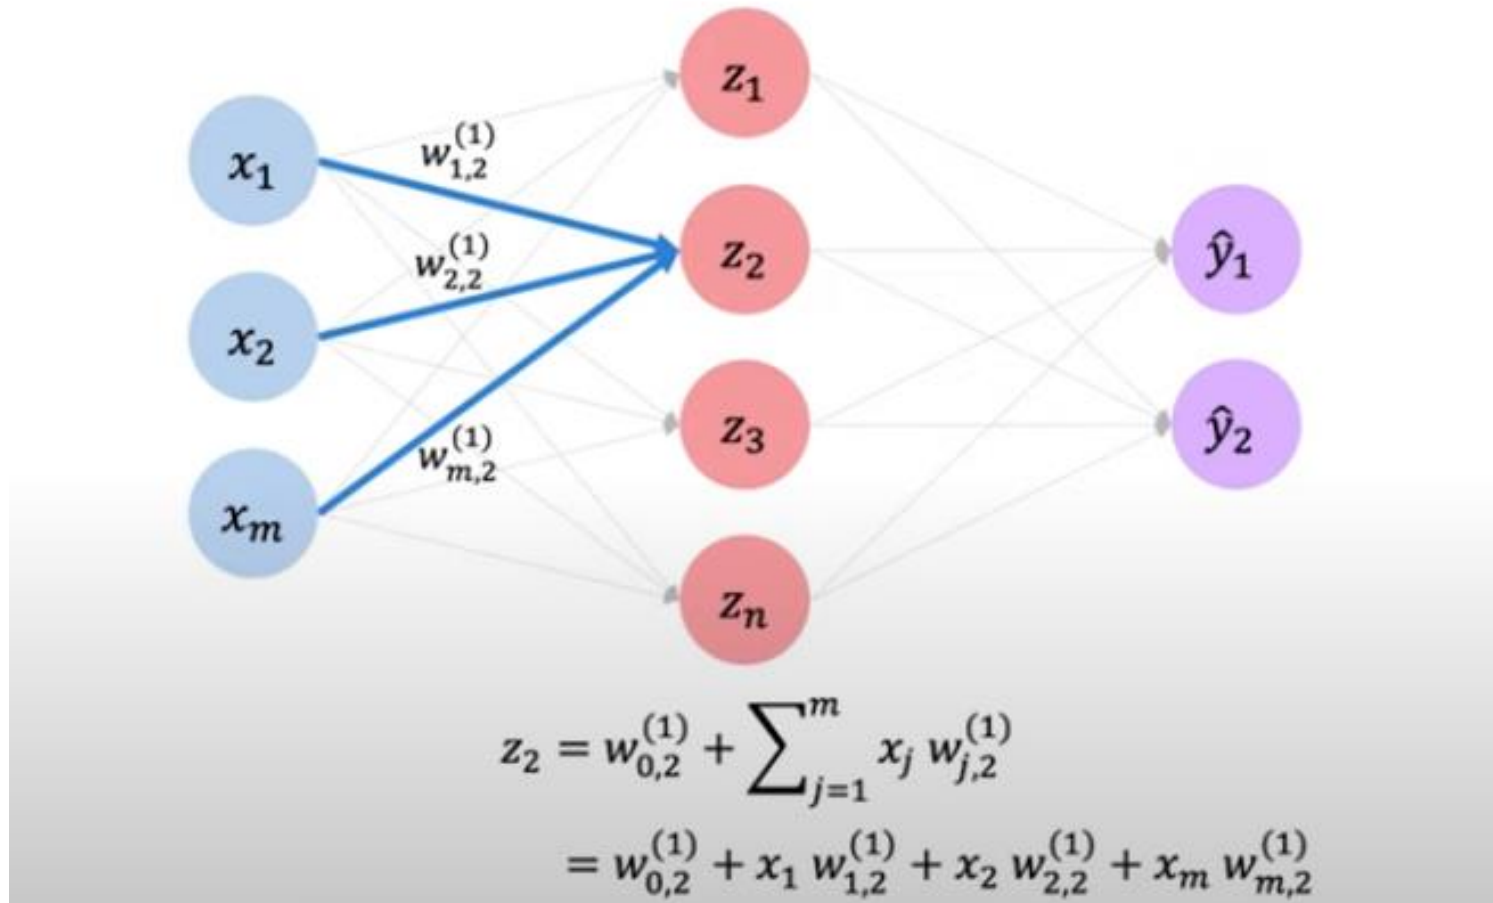


1. The layer in between the input and output layer is called the hidden layer
2. The state of the hidden layers are typically unobserved or hidden.

# The single layer ANN



**NUS**  
National University  
of Singapore



An example of computing the hidden layer second neuron ( $z_2$ )



# How do you train the ANN?

ANN are trained using gradient descent and require that you use a loss function when designing and configuring your ANN model. Note: Loss function are sometimes known as Objective Function or Cost Function or Error Function.

Loss function plays a critical role in training ANN as it provide a means to calculate the model error. Typically one would like to minimise the error hence minimising the loss function

Mean squared error (MSE) is the main type of loss functions to use when training neural network models.


# Loss Function and Optimisation

$$W^* = \underset{W}{\operatorname{argmin}} J(W)$$

Remember:  
*Our loss is a function of  
the network weights!*



**NUS**  
National University  
of Singapore



Our loss function  $J(W)$  is to optimise all the weights. Loss function here is minimising the error of the network, and to do so, is to find the lowest points in this landscape. We do that by computing the gradient of the loss function  $J(W)$  with respect to the weight.

# Loss Function using Mean Squared Error (MSE)



This is calculated as the average of the squared differences between the predicted and actual values. The result is always positive regardless of the sign of the predicted and actual values and a perfect value is 0.0. The loss value is minimized



The Python function below provides a pseudocode-like working implementation of a function for calculating the mean squared error for a list of actual and a list of predicted real-valued quantities.

```
# calculate mean squared error
def mean_squared_error(actual, predicted):
    sum_square_error = 0.0
    for i in range(len(actual)):
        sum_square_error += (actual[i] - predicted[i])**2.0
    mean_square_error = 1.0 / len(actual) * sum_square_error
    return mean_square_error
```




# Loss Function and Optimisation



**NUS**  
National University  
of Singapore

Compute gradient,  $\frac{\partial J(w)}{\partial w}$




How do we start? First pick a point for  $w_0$  and  $w_1$ ; then find the lowest point by computing the gradient of that point. Take the negative gradient, as we want to find the lowest point in the landscape. Lowest point here is the minimum error.

# Loss Function and Optimisation



**NUS**  
National University  
of Singapore

Take small step in opposite direction of gradient




# Gradient Descent



**NUS**  
National University  
of Singapore

Repeat until convergence



Repeat until it converge to the lowest point.

# Gradient Descent algorithm




**NUS**  
National University  
of Singapore

## Algorithm

1. Initialize weights randomly  $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3.   Compute gradient,  $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
4.   Update weights,  $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
5. Return weights

The small learning rate value lead to a small step.

# Computing Gradient - Backpropagation




*How does a small change in one weight (ex.  $w_2$ ) affect the final loss  $J(W)$ ?*

Start with a simple neural networks that have one input, one hidden layer with one neuron and one output. Compute the gradient of the loss function  $J(W)$




# Computing Gradient - Backpropagation



$$\frac{\partial J(W)}{\partial w_2} = \underbrace{\frac{\partial J(W)}{\partial \hat{y}}}_{\text{purple}} * \underbrace{\frac{\partial \hat{y}}{\partial w_2}}_{\text{red}}$$

# Computing Gradient - Backpropagation




$$\frac{\partial J(W)}{\partial w_1} = \frac{\partial J(W)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial w_1}$$

↑  
Apply chain rule!
↑  
Apply chain rule!

Question: How do we compute the gradient of the loss function with respect to  $w_1$ . Apply the chain rule.

# Computing Gradient - Backpropagation



$$\frac{\partial J(W)}{\partial w_1} = \underbrace{\frac{\partial J(W)}{\partial \hat{y}}}_{\text{purple}} * \underbrace{\frac{\partial \hat{y}}{\partial z_1}}_{\text{red}} * \underbrace{\frac{\partial z_1}{\partial w_1}}_{\text{blue}}$$

Small change in the weight will affect our loss.

Repeat this chain rule steps for all the weights in the subsequence layers (for deep neural network you have many layers)

# Loss Functions can be difficult to optimise

**Remember:**

Optimization through gradient descent

$$\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$$

# How best to optimise the loss function?

## Remember:

Optimization through gradient descent

$$\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$$

How can we set the  
learning rate?

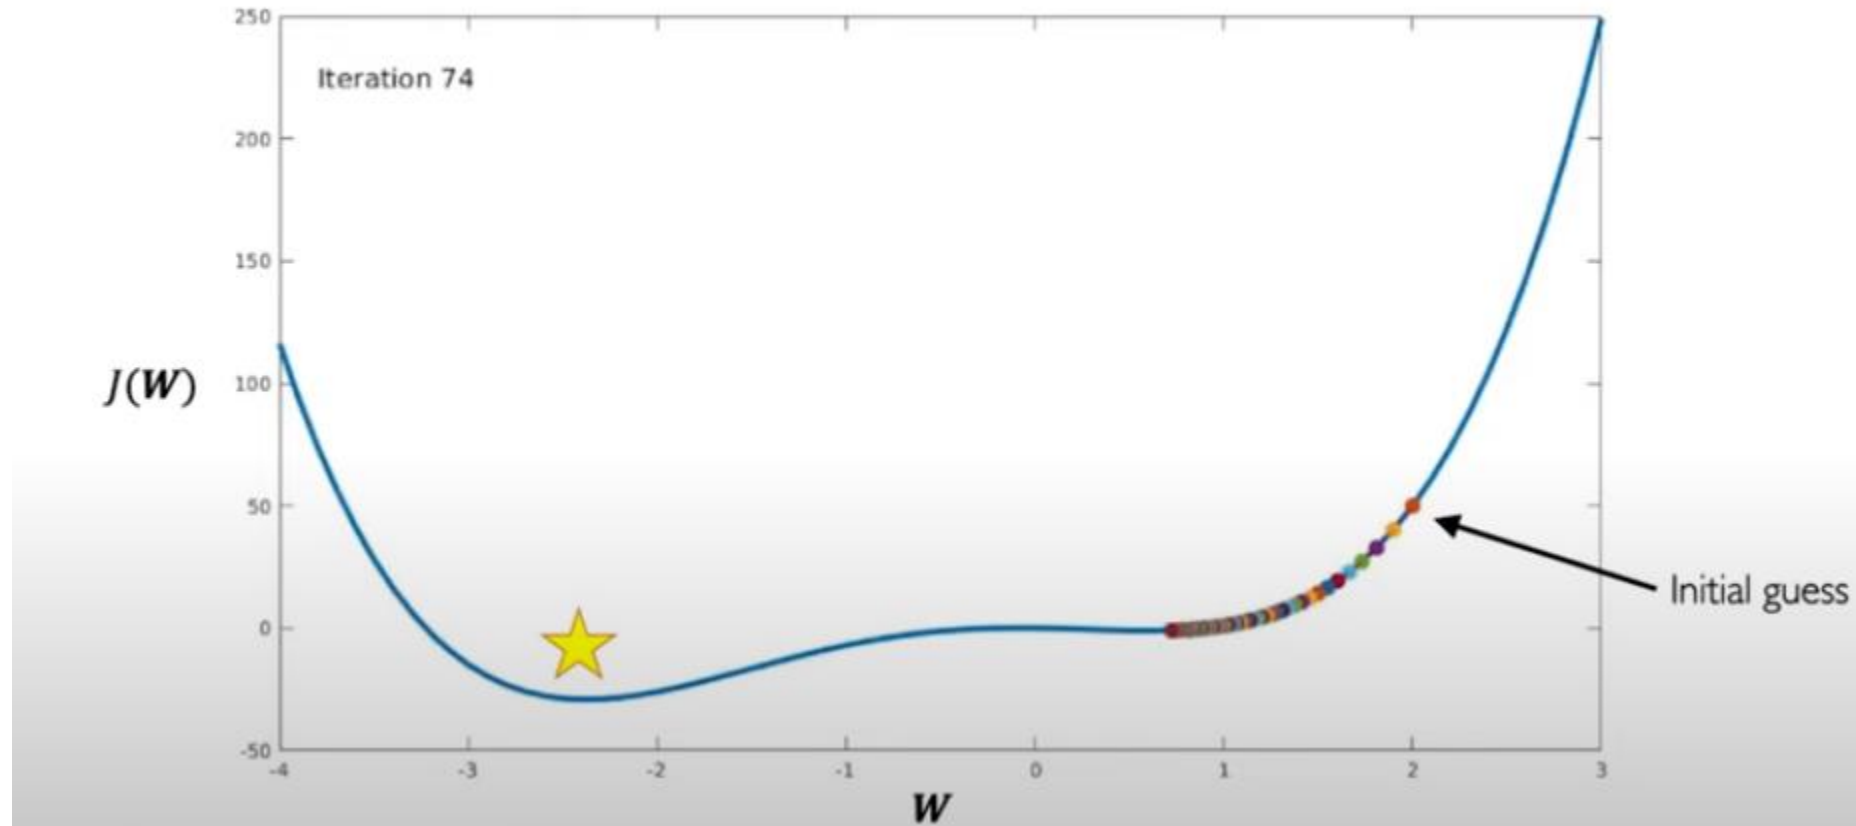
Eta ( $\eta$ ) is the learning rate and is the step size of the moment or the update to the weight.

# Setting the learning rate



**NUS**  
National University  
of Singapore

*Small learning rate converges slowly and gets stuck in false local minima*




If the learning rate is too small it may stuck at local minimum or learn too slowly

# Setting the learning rate



**NUS**  
National University  
of Singapore

*Large learning rates overshoot, become unstable and diverge*



Large learning rate may cause overshoot and cannot converge.




# Setting the learning rate



**NUS**  
National University  
of Singapore

*Stable learning rates converge smoothly and avoid local minima*



Selecting the right learning rate and having a reasonable starting point.  
What is a typical learning rate and what is a good starting point?

# Adaptive learning rates

- Learning rates are no longer fixed
- Can be made larger or smaller depending on:
  - how large gradient is
  - how fast learning is happening
  - size of particular weights
  - etc...

# Multiple Output Perceptron also known as Multi-Layer Perceptron (MLPs)




**NUS**  
National University  
of Singapore

## *Generalization of the single-layer perceptron*

Consists of:

- An input layer
- One or more hidden layers of computation nodes
- An output layer of computation nodes

Architectural graph of a multilayer perceptron with two hidden layers:



# Assignment on ANN



**NUS**  
National University  
of Singapore

a.) Draw and label an ANN with one input layer, one hidden layer and one output layer; the number of neurons in the input layer is 2; and number of neuron in the hidden layer is 3. The output layer will have only 1 neuron. The activation function for the hidden neuron is sigmoid function. The activation function for output layer is any linear function. The input layer neurons have no activation function.

b.) List all the equations needed to calculate this ANN for both forward pass and backward pass.

c.) Demonstrate the training and learning of this ANN in doing classification or function approximation task. Note: To demonstrate the working of this simple ANN, you can generate your own data points.

d.) Show all diagrams, graphs and the results of the classification task.



**NUS**  
National University  
of Singapore

Thank you for your attention  
Q&A during tutorial time